



# Ballerina

**An open-source programming language for the cloud that makes it easier to use, combine, and create network services.**

June 2021

# The background to Ballerina



## 15 years and 1000s of customers taught us that ...

- All enterprises need lots of integration to innovate digitally
- Integration remains hard, time consuming and expensive
  - This is applicable to all technology vendors, not just us
    - See [MuleSoft's 2021 Connectivity Benchmark Report](#)
- Cloud native engineering requires integration systems to be simply code that runs in containers
  - The days of big servers running middleware as a central service are over
- Integration needs to be simpler, less time consuming and cheaper





**Integration is programming, but...**

A **visual representation of integration logic** is important to communicate with business users.

**Domain specific languages (DSLs)** have dominated because they provide the right abstractions for integration programming, albeit with limitations when it comes to “regular code” parts of the problem.

**Integration programming** has lost software engineering **best practices** because it lives in a closed universe.

# The Ballerina project





## The Ballerina project

- Started in 2016
- More than 300 person years of investment to date
- The language, its libraries and tools are all open source under Apache License

# The Ballerina Programming Language

- Addresses most use cases of DSLs & scripting languages, but with the scalability and robustness of application languages
- Designed for mass usage - not an “elite” language
- Designed to support a graphical view
- Designed for the cloud
- “Batteries included” - comes with support for:
  - Package system
  - Structured documentation
  - Testing
  - Lots of libraries for network data, messaging & communication protocols





# Features of Ballerina



# Data oriented

```
// closed type
type Coord record {|
    float x;
    float y;
|};
```

```
Coord coord = { x: 1.0, y: 2.0 };
```

```
// nothing to do
json j = coord;
```

```
// If coord is open:
type Coord record {
    float x;
    float y;
};
```

```
// usually happens automatically
json j = coord.toJson();
```

- Object-orientation bundles data with code: wrong approach for network interaction
- Ballerina emphasizes plain data - data that is independent of any code used to process the data
- Ballerina provides objects for internal interfaces, but is not object-oriented
- Ballerina's plain data maps straightforwardly to and from JSON



# Powerful features for working with data

```
type Employee record {  
    string firstName;  
    string lastName;  
    decimal salary;  
};  
  
Employee[] employees = [  
    // ...  
];  
  
Employee[] sorted =  
    from var e in employees  
    order by e.lastName ascending,  
            e.firstName ascending  
    select e;
```

- Language integrated query with SQL-like syntax
- Table data type - work with relational and tabular data
- XML support - integrates functionality similar to XQuery
- Decimal data type - numbers designed for the needs of business

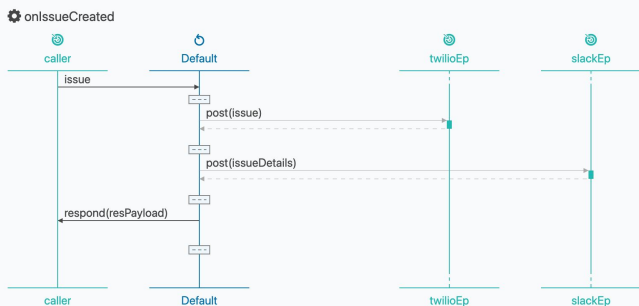


## **A Ballerina type system is structural and works as a schema language**

- Structural, not nominal typing provides looser coupling
- Type system supports open structures: say as much or as little about the structure as you need to
- Static typing, but some things are checked at runtime in order to make the type system less complicated and more flexible
- Works for describing both the operations the program performs and data on the wire
- Works as a schema for network data as well as a type system - eliminates “data binding” problem particularly for JSON

# Text and graphical syntax parity

```
remote function onIssuesCreated
  (webhook:IssuesEvent event)
    returns error? {
      webhook:Issue issueInfo = event.issue;
      ...
    }
```



- A function has equivalent representations as both a textual syntax and a sequence diagram
- The sequence diagram provides insight into a function's network interactions and use of concurrency
- Horizontal line for messages sent
- Only possible because it has been designed into the language from the start
  - Design of function-level concurrency features
  - Language has network abstractions



# Inherently concurrent

```
function post(string message) {  
  worker T {  
    var r = twitter -> tweet(message);  
    r -> function;  
  }  
  
  worker L {  
    var r = linkedIn -> post(message);  
    r -> function;  
  }  
  
  var twitterResp = <- T;  
  io:println("Twitter: ", twitterResp);  
  
  var linkedInResp = <- L;  
  io:println("LinkedIn: ", linkedInResp);  
}
```

- Main concurrency concept is a "strand": a thread in the concurrency sense (similar to goroutines in Go)
- A strand is cheap: does not require an OS thread
- A function can have named "workers" that each run on a new strand concurrently with the function's default worker
- Strands can be scheduled on separate OS threads to provide parallelism
- Provides the advantages of async functions with simpler programming model



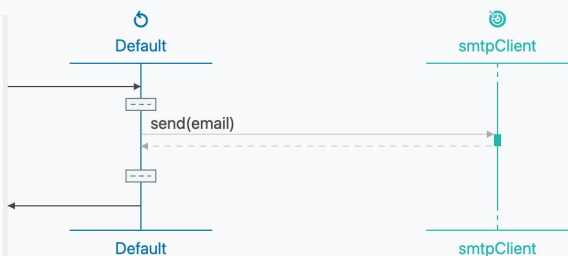
# Consumes network services

```
import ballerina/email;

function main() returns error? {
    email:SmtpClient sc
        = check new("smtp.example.com",
                    "user123@example.com",
                    "passwd123");
    check sc -> sendEmail({
        to: "contact@ballerina.io",
        subject: "Ballerina"
        body: "I love Ballerina!"
    });
}
```

- Key enabler for sequence diagram view of network interactions
- Outbound network interactions represented by client objects
- Client objects have remote methods that represent outbound interactions with a remote system
- Distinct syntax for calls remote method
- Syntax restrictions make it possible to create a sequence diagram for any function

f main



# Produces network services

```
import ballerina/http;

service / on new http:Listener(9090) {
    resource function get hello(string
        name) returns string {
        return "Hello, " + name;
    }
}
```

- Application defines service objects and attaches them to Listeners
- Libraries provide protocol-specific Listeners, which receive network input and dispatch to service objects
- Service objects support two interface styles
  - remote methods, named by verbs, support RPC style
  - resources, named by method (e.g. GET) + noun, support RESTful style (used for HTTP and GraphQL)
- Types of service objects methods can used to generate interface descriptions e.g. OpenAPI, GraphQL
- Annotations on service objects enable easy cloud deployment





## Concurrency safety

- Ballerina allows strands to share mutable state in order to provide a familiar programming model
- Combination of concurrency and shared mutable state creates potential for data races (silently giving an incorrect result)
- For function workers, Ballerina avoids races by cooperatively multitasking all strands onto a single thread
- Type system has features that makes it possible to determine when services have locked enough to be able to safely use multiple threads to handle incoming requests in parallel
- Does not give massive parallelism, but enough to make effective use of common cloud instance types

# Transactions

- Makes it easier to write Ballerina programs that use transactions
- Not transactional memory
- Language support for delimiting transactions
- Ballerina runtime includes transaction manager
- Composes with network interaction features to support distributed transactions

# Error Handling

- Error handling approach has pervasive impact on language design and usage
- Errors are normal when you are dealing with a network
- Exceptions are the wrong approach for dealing with normal errors
  - Control flow is implicit
  - Code is harder to understand and maintain
- Trend in modern application/system languages is for error control flow to be explicit: Go, Rust, Swift
- Scripting languages typically use exceptions
- Ballerina uses error data type with explicit error control flow

# Error handling example

```
configurable string host = ?;
configurable string username = ?;
configurable string password = ?;

public function main() {
    error? err = sendEmail(to = "contact@ballerina.io", subject = "Ballerina", body = "I love Ballerina");
    if err is error {
        io:println(`Error sending email: ${err.message()}`);
    } else {
        io:println("Email sent!");
    }
}

function sendEmail(string to, string subject, string body) returns error? {
    email:SmtpClient smtpClient = check new (host, username, password);
    check smtpClient->sendMessage({to, subject, body});
}
```



## **Ballerina is meant to be familiar**

- Popular C-family languages (C, C++, Java, JavaScript, C#, TypeScript) have a lot in common
- Ballerina leverages this: there's a subset that is enough to get started and will feel very familiar to any programmer with experience of one of these C-family languages
- Ballerina provides better ways to do things, but also familiar ways to work

## Ballerina offers not just the language, but the full platform

- VSCode plugin
  - Source and graphical editing
  - Debugging
- Tools for working with OpenAPI, GraphQL schemas, gRPC schemas
- Ballerina Central
  - Module sharing platform
- Integration to [Choreo by WSO2](#) for observability, CI/CD and DevOps



# Current status

## Ballerina 1.0

Released in 2019

## Ballerina 2021 (Swan Lake)

Major new version

Public Beta launched June 2nd  
2021

Extended review / verification  
period before GA



# Ballerina implementations

## jBallerina

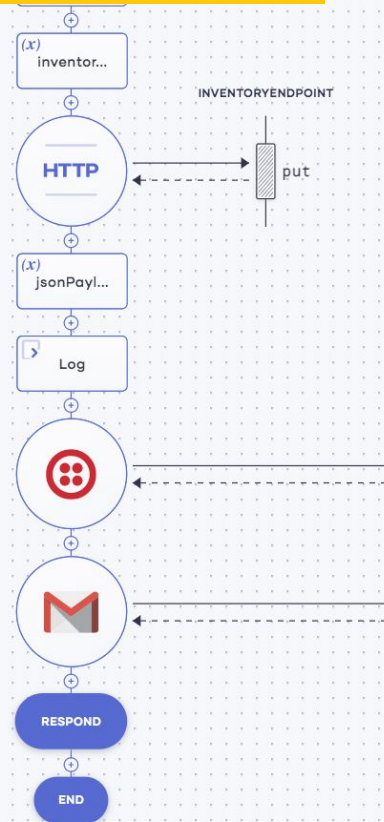
- Toolchain written using Java
- Compiles to Java bytecodes and runs on a JVM
- Provides Java interoperability
- Available now

## nBallerina

- Cross compilation to native binaries via LLVM
- Toolchain will be shared initially (compiler front-end still in Java) but fully bootstrapped soon
- Provides a C FFI
- ETA: initial release end of 2021



# Summary



```
ballerina/http;

on new http:Listener(8090)
source function post orders(

    json orderPayload = checkpanic
    var orders = checkpanic or
    var contactNo = checkpanic
    var email = checkpanic or
    var price = checkpanic or
    var inventoryURL = "https:

http:Client inventoryEndpo
http:Response inventoryRes
var jsonPayload = checkpan

log:print(jsonPayload.toJs

twilio:Client twilioEndpoi
    accountSID: "ACd0fd3df
    authToken: "0744c6a75d
    xAuthyKey: ""
```

- Ballerina is a modern, industrial grade programming language optimized for writing integrations in a cloud native environment
- Type system is designed to make network data processing easier
- First class network services along with functions/objects
- Fully open source and developed openly
- Sponsored by WS02



# Thanks!



[wso2.com](http://wso2.com)

